

# Data Engineering

- Data Processing
- Batching
- Spark

Data Engineering Course  
Lucas Rosa  
2024

# Data Processing

## Introduction

### Lucas Porto Rosa

Brazilian, 36 years

Principal Data Engineer at HBDC (Metys & Hugo Boss)

LinkedIn: <https://www.linkedin.com/in/lucprosa/>



### Personal

Moved to Portugal in Oct 2021 with wife and dog

Gaúcho, gremista

Crafter beer, travel, guitar player, etc.

### Work Experience

> 10 years of experience working with data as DBA, Business Intelligence analyst and Data Engineer

### Education

System Development Analysis

MBA Data Science

Tech Certifications



# Data Processing

## Introduction

Data Engineer experience?

Spark?

SQL?

Python?



# Data Processing

## Introduction

### Development Environment

Google Colab - <https://colab.research.google.com/>

GitHub - <https://github.com/lucprosa/dataeng-basic-course/>

Dataproc - <https://cloud.google.com/>

# Data Processing

## Agenda

	<b>15 -Data Processing</b>  Tutor: Lucas Rosa Horário: 19h - 23h	<b>16 -Data Processing</b>  Tutor: Lucas Rosa Horário: 9h - 18h	17
	<b>22 -Data Processing</b>  Tutor: Lucas Rosa Horário: 19h - 23h	<b>23 -Real-Time Data (Streaming)</b> Tutor: Lucas Rosa Horário: 9h - 18h	24
	29	<b>30 -Real-Time Data (Streaming)</b> Tutor: Lucas Rosa Horário: 9h - 18h	

- Day 15
  - Data Processing / Batching Introduction
  - Spark Introduction, components
  - Hands-On (Google Colab + Dataproc)
  - Spark architecture
  - Data Solutions / Alternatives to Spark
- Day 16
  - Spark common Issues
  - Hands-On (Google Colab + Dataproc)
  - Concepts about ETL/Medallion Architecture
  - Technical challenge
- Day 17
  - Tech challenges (continuation)
  - Doubts/Question



# Data Processing

## Index

- Batch Processing
- Apache Spark
  - Introduction & history
  - MapReduce vs Spark, Hadoop
  - Spark Components
  - Spark Architecture
  - Common issues / Performance
  - Code examples
  - Hands-on
- ETL, Lakehouse, Medallion Architecture
- Technical Challenge

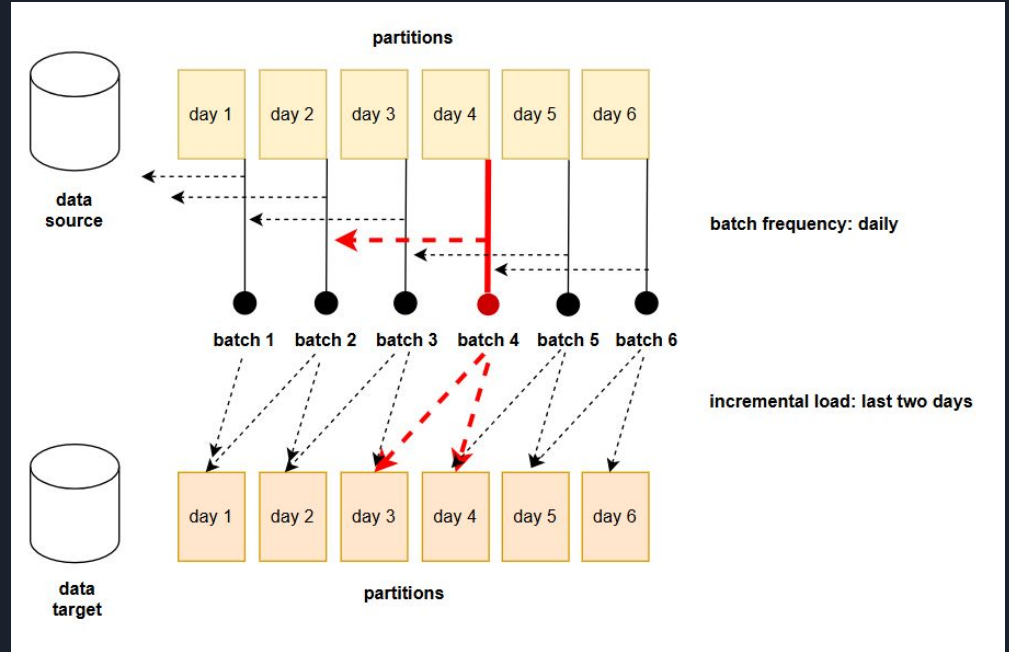
# Data Processing Batching



# Data Processing

## Batching

- Batch jobs
- Data is collected, stored and processes in batches
- Jobs are scheduled / batch frequency
- Full and incremental loads

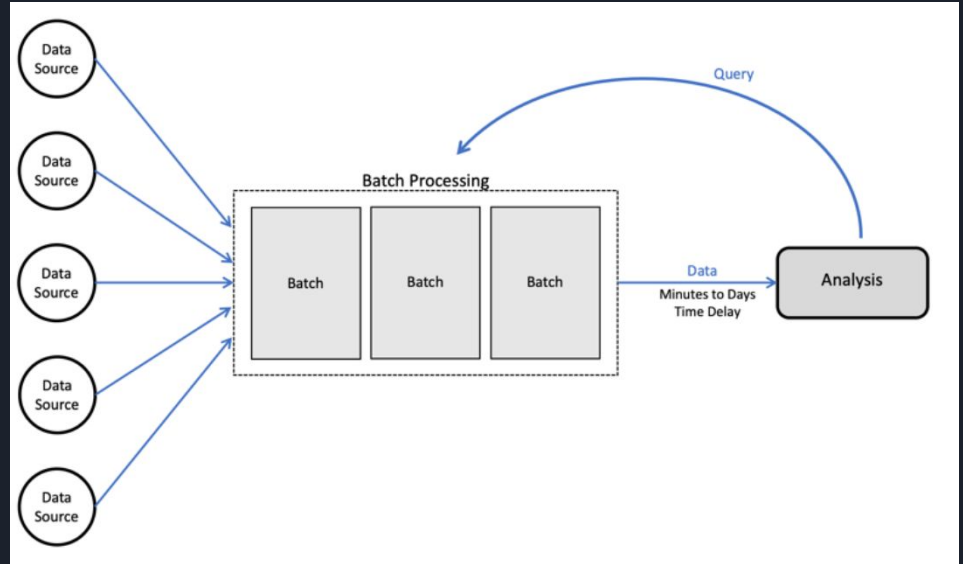




# Data Processing

## Batching

- Process data in batches / chunks
- **Data Volume:** Large amount of data
- **Data Latency:** High latency (hourly, daily, weekly, monthly)
- **Cost:** Low cost (comparison to streaming)





# Data Processing

## Batching

### USE CASES

- Data Integration, data consolidation
- ETL/ELT jobs
- Data Quality
- Data Archiving
- Backups
- Data Mining

### BENEFITS

- Data Analytics
- Dashboarding
- Reports
- Machine Learning
- Monitoring KPIs
- Business decision-making
- Security alerts
- Data Quality checks
- Data transformations & enrichments

# Apache Spark





# Data Processing

Apache Spark

- Apache Spark
- <https://spark.apache.org/>




# Data Processing


## Apache Spark

- Unified engine for large-scale data analytics
- Distributed and parallel computing, in-memory, fault tolerant, etc
- Open Source (<https://www.apache.org/>)
- Developed in 2009 (UC Berkeley) to replace MapReduce computing paradigm
- Spark cluster components (driver/master node, workers, executors)
- RDDs (resilient distributed dataset)

Distributed/Parallel Processing



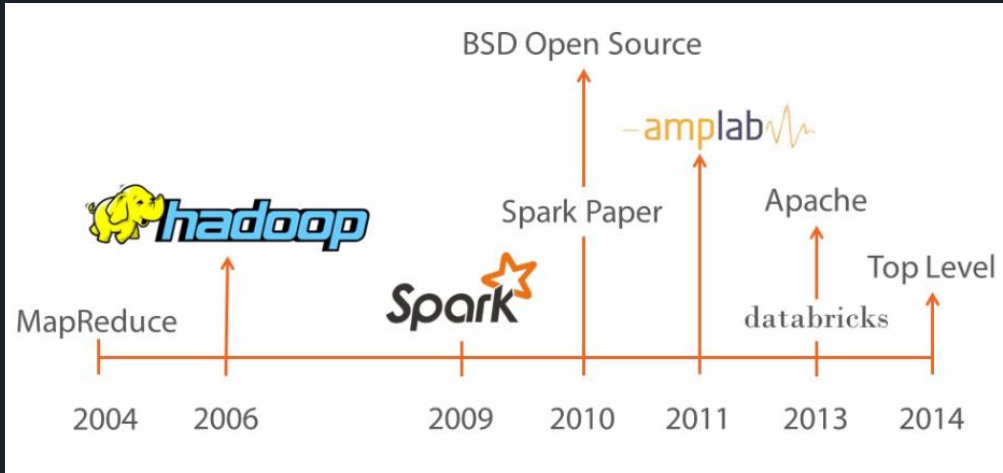
Spark architecture



# Data Processing

## Apache Spark

- Big Data problems
  - How to store?
  - How to process?


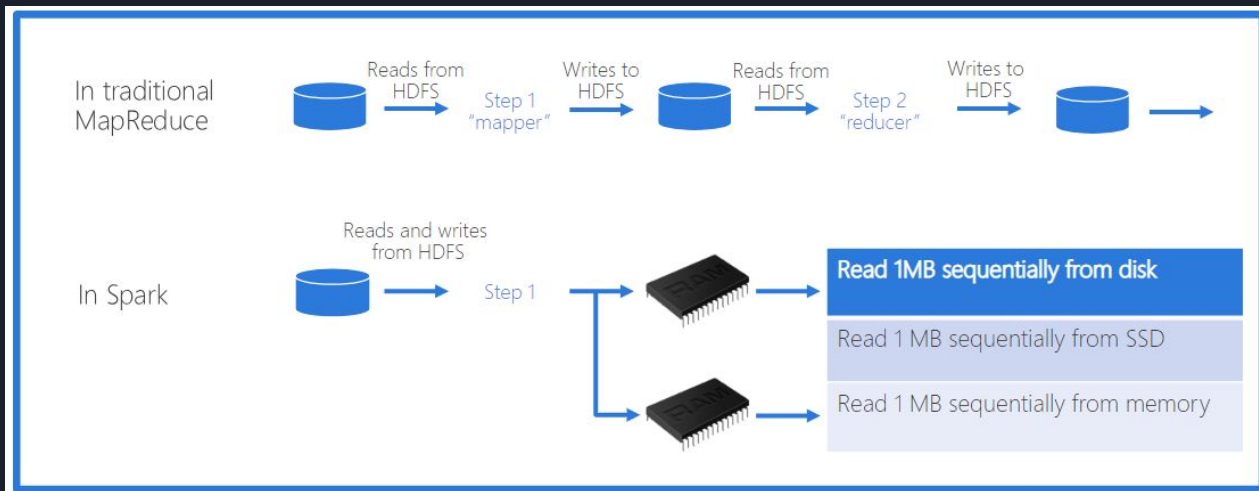


- 2002 - Hadoop - Apache Nutch
- 2003 - Google - GFS (Google File System)
- 2004 - Google - MapReduce
  - Big data processing model
  - Java focused
  - Map, Shuffle, Partition, Reduce
  - Read/write intensive
- 2004 - Hadoop - Apache Nutch
  - GFS + MapReduce
- 2006 - Hadoop - Yahoo + Apache Nutch
  - HDFS + MapReduce
- 2008 - Apache Hadoop
- 2009 - Spark - Research at UC Berkeley AmpLabs
- 2010 - Spark - First paper
- 2013 - Spark - Apache Software Foundation
- 2014 - Spark 1.0

# Data Processing

Apache Spark

- MapReduce vs Spark



<https://amplab.cs.berkeley.edu/projects/spark-lightning-fast-cluster-computing/>

⌋



# Data Processing

Apache Spark

- MapReduce vs Spark

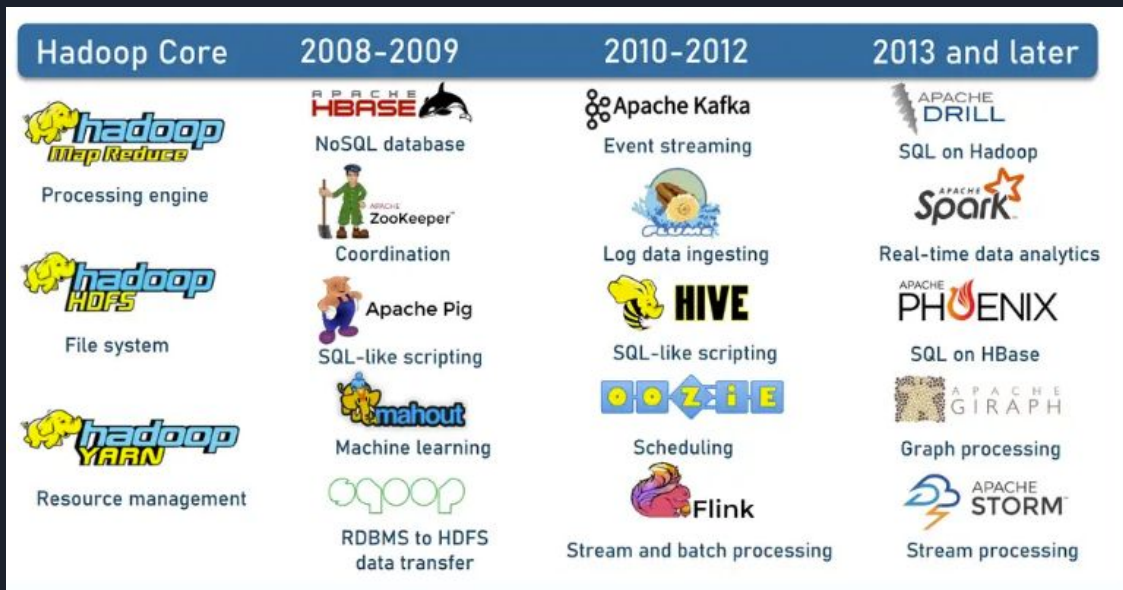
Criteria	Hadoop	Spark
Real-time Data Processing	Primarily for batch processing, not optimized for real-time tasks.	Well-suited for real-time or near-real-time processing due to in-memory speed.
Accessing Data Randomly in Memory	Reads and writes data to/from disk, less efficient for random memory access.	Designed for in-memory processing, allowing efficient random data access.
Iterative and Interactive Operations	Writes intermediate results to disk, and can be slow for iterative or interactive tasks.	Optimized for iterative and interactive operations, keeping data in memory.



# Data Processing

## Apache Spark

- Hadoop Ecosystem timeline



# Data Processing

Apache Spark

Spark - How/where to use it?

Local machine,  
standalone, K8




SaaS / PaaS



Snowpark



Hadoop cluster on  
Cloud



Hadoop cluster on  
Prem



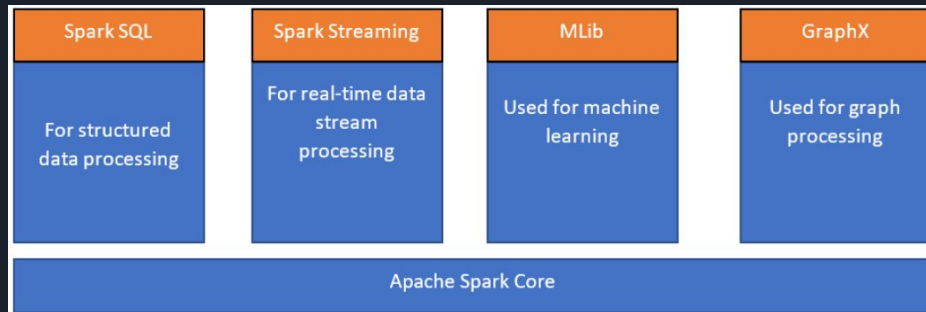

# Apache Spark Architecture



# Data Processing

## Apache Spark

- APIs
- Components
- Architecture



Reference: <https://spark.apache.org/docs/latest/cluster-overview.html>



# Data Processing

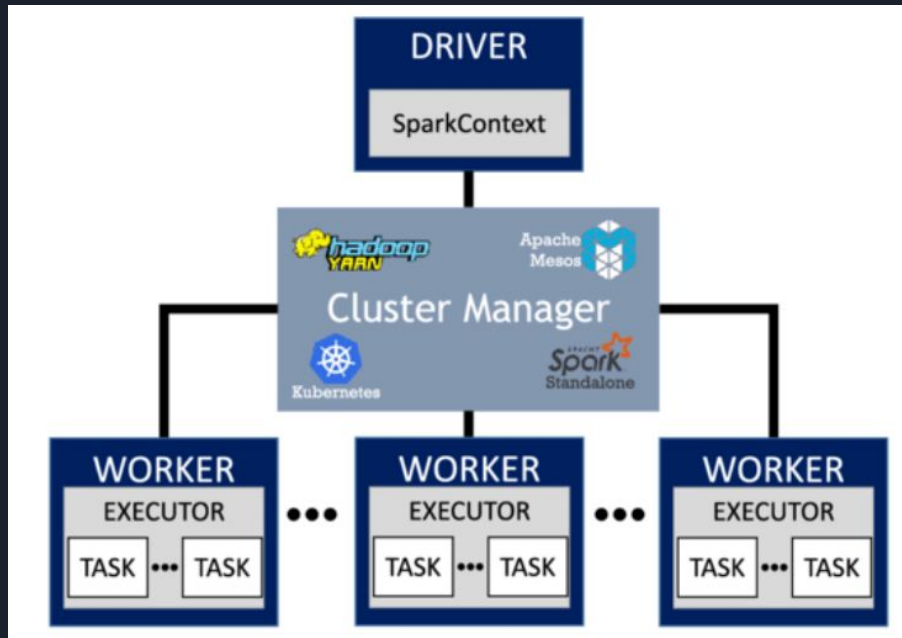
Apache Spark

- **Spark APIs (languages)**
  - Java, Scala, Python, R
- **Spark APIs**
  - **SQL API**
  - Dataset API (only Scala and Java)
  - **Dataframe API**
  - Pandas API
  - MLib for machine learning
  - GraphX for graph processing
  - **Structure Streaming** for stream processing
  - Spark Connect API (> 3.4)
- **Spark CLI**
  - spark-shell
  - spark-submit
  - **pyspark**
  - sparkR
  - spark-sql

# Data Processing

## Apache Spark


- Spark Components
  - Driver
  - Workers
  - Cluster Manager
    - YARN
    - MESOS
    - Kubernetes



# Data Processing

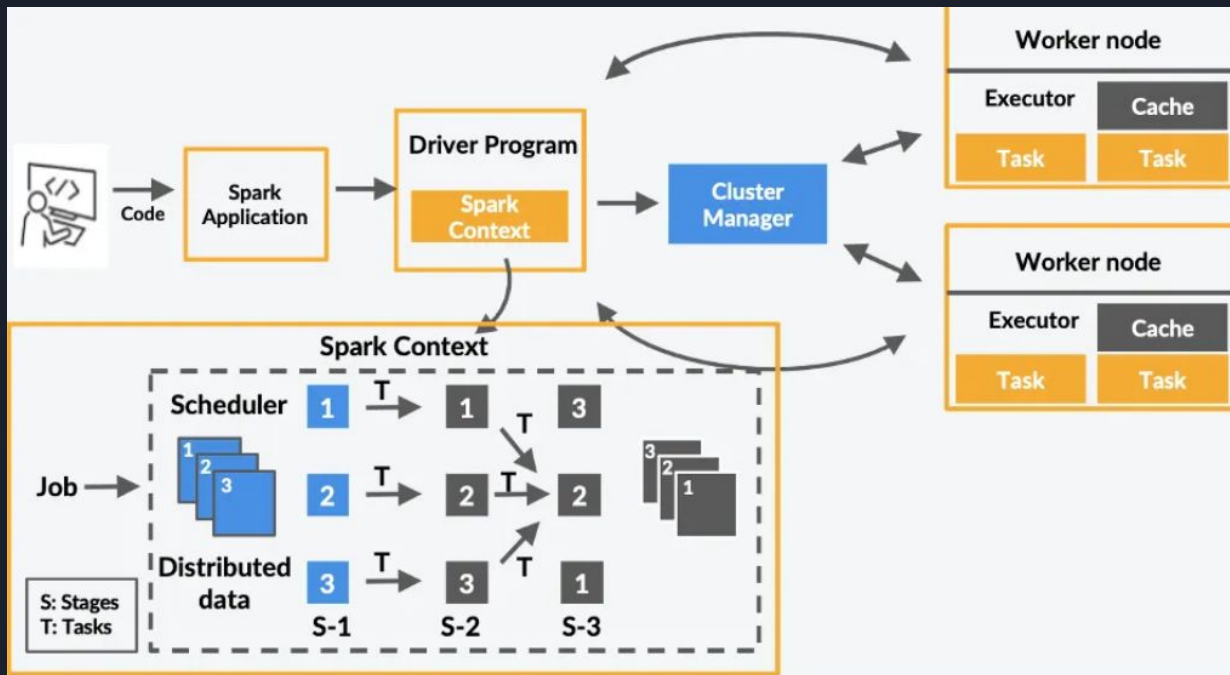
## Apache Spark

- **Workers** - Machines in the cluster
- **Driver** - Central control for Spark application (main method)
- **Cluster Manager** - Launches executors and allocate and manage resources
- **Executors** - Processes running tasks on workers (can take one partition at the time)
- **Cores** - CPU cores allocated to executors
- **Cache** - Memory or disk caching in workers
- **Tasks** - Spark commands sent by Driver to Executors
- **Stages** - Contains a quantity of tasks
- **Partitions** - Logical chunk of data in a large distributed dataset (128MB)



# Data Processing

Apache Spark



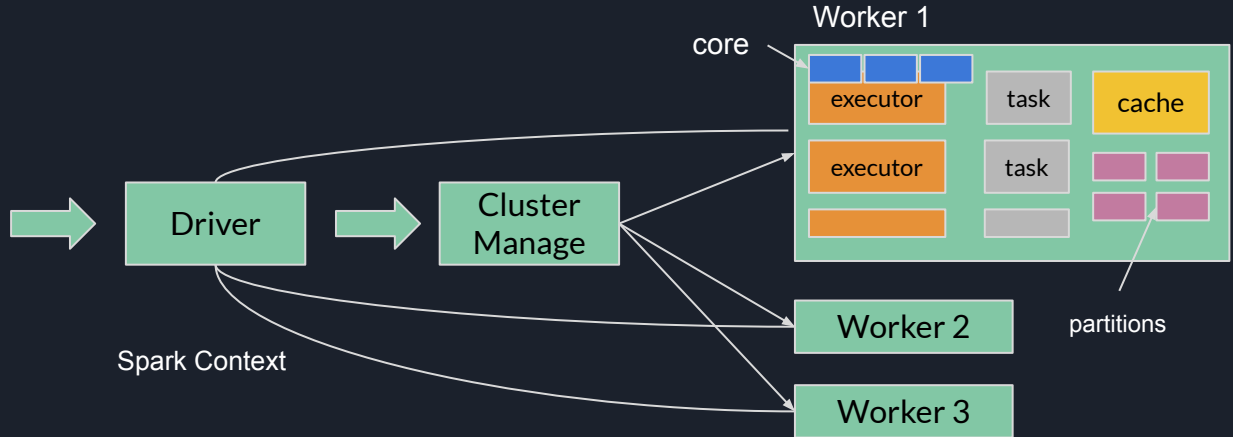


# Data Processing

Apache Spark

## CODE:

line 1 - read csv 1 from path 1 -> df1  
line 2 - read csv 2 from path 2 -> df2  
line 3 - add new column to df1 -> df1  
line 4 - add new column to df2 -> df2  
line 5 - join df1 with df2 -> df3  
line 6 - aggregate data -> df4  
line 7 - write df4 to the lake



- code is splitted in tasks and spread among the workers by the driver
- RDDs will be splitted among the partitions in the workers
- executors run one task per core
- single task will operate on a single partition

- parallelism
- distributed workload
- fault-tolerance (operation lineage)

## To evaluate Spark parallel tasks and performance

- How many workers does the cluster have?
- How much Memory and CPU each worker have?
- Define memory allocation and CPU for each executor and driver
- Define shuffle partitions



# Data Processing

Apache Spark

- **RDD**

Resilient distributed dataset (Fault-tolerant collection of elements that can be operated on in parallel)

Ex: read data from parquet (1GB) → rdd (contains many partitions distributed across the cluster)

- **TRANSFORMATIONS**

Operations that return another RDD/dataset as output

Ex: join, groupBy, filter

- **ACTIONS**

Operations that return a value to the driver after running computation on the RDD/dataset

Ex: write, count, show, collect

- **DAG**

Logical execution plan for a job. Sequence of operations.

Ex: T1 -> T2 -> T3 -> T4 -> A1



# Data Processing

Apache Spark

- **DATASETS**

Distributed collection of data (strong typing, ability to use powerful lambda functions)  
RDD's benefits + Spark SQL's optimized execution engine

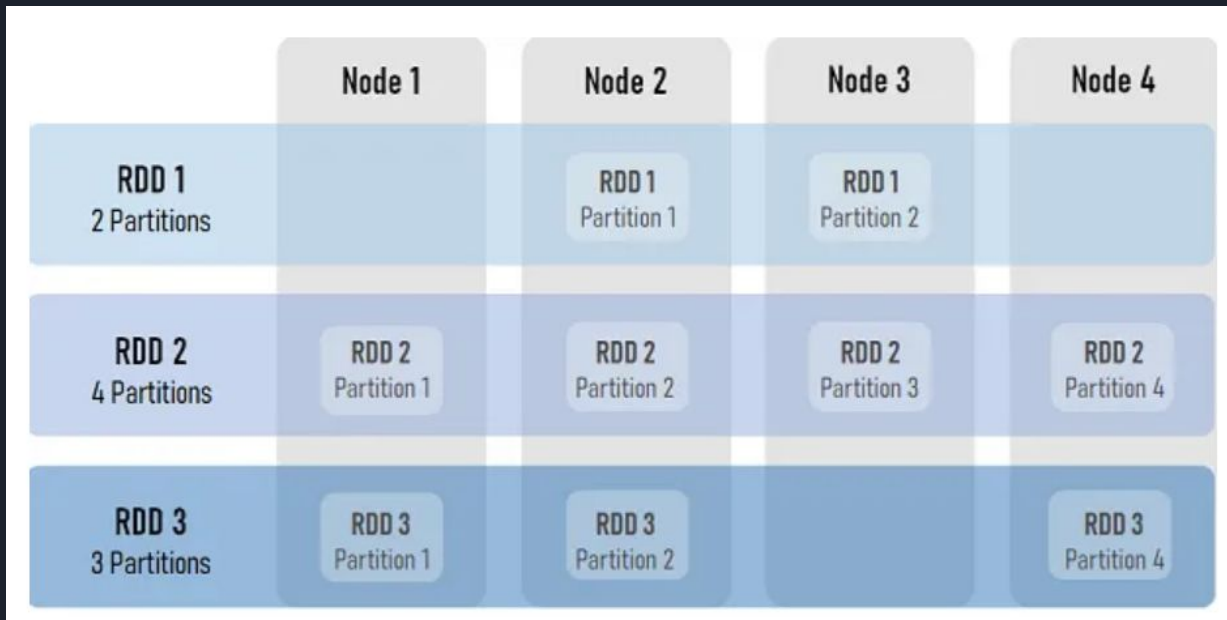
- **DATAFRAMES**

Dataset organized into named columns. Equivalent to a table in a relational database  
RDD's benefits + Spark SQL's optimized execution engine

# Data Processing

Apache Spark

RDDs x Partitions x Workers(Nodes)



# Spark Hands-On





# Data Processing

Apache Spark

## HANDS-ON

<https://github.com/lucprosa/dataeng-basic-course/tree/main>

spark/examples/00-setup.ipynb	Installation, Spark Session, Spark Context
spark/examples/01-rdds.ipynb	RDDs
spark/examples/02-dataframes.ipynb	DataFrame
spark/examples/03-sql.ipynb	SQL, Temp Views, Spark Catalog
spark/examples/04-joins.ipynb	Joins
spark/examples/05-aggregations.ipynb	Aggregations
spark/examples/06-write_partitioning.ipynb	Writing operation, Write Mode, partitionBy
spark/examples/07-udf.ipynb	User-defined functions
spark/examples/09-windows-function.ipynb	Windows Function
spark/examples/10-misc_performance.ipynb	Cache, Persist, broadcast join, repartition/coalesce, explain

# Data Processing

Apache Spark

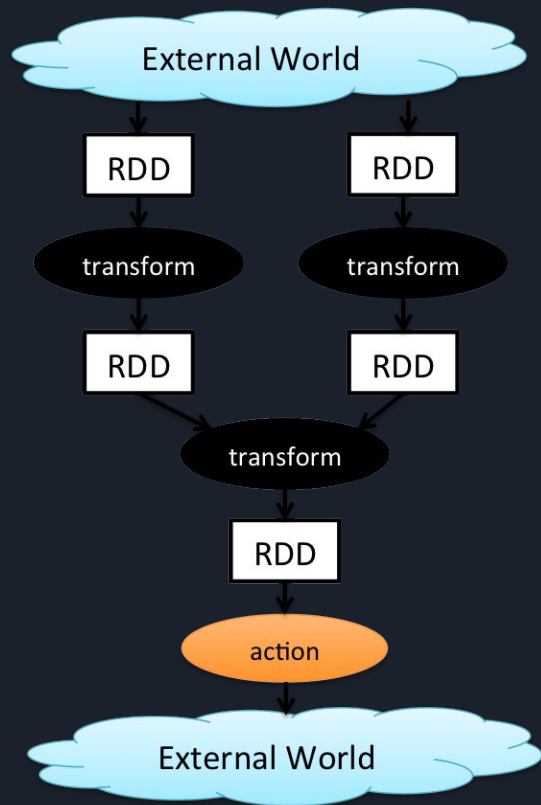
- Lazy Evaluation - Action vs Transformation
- DAG - Directed Acyclic Graph
- Wide & Narrow Transformations
- Data Shuffling

## Transformations

- Wide (multiple partitions) - groupBy, join, distinct...
- Narrow (single partition) - map, filter, union...

## Actions

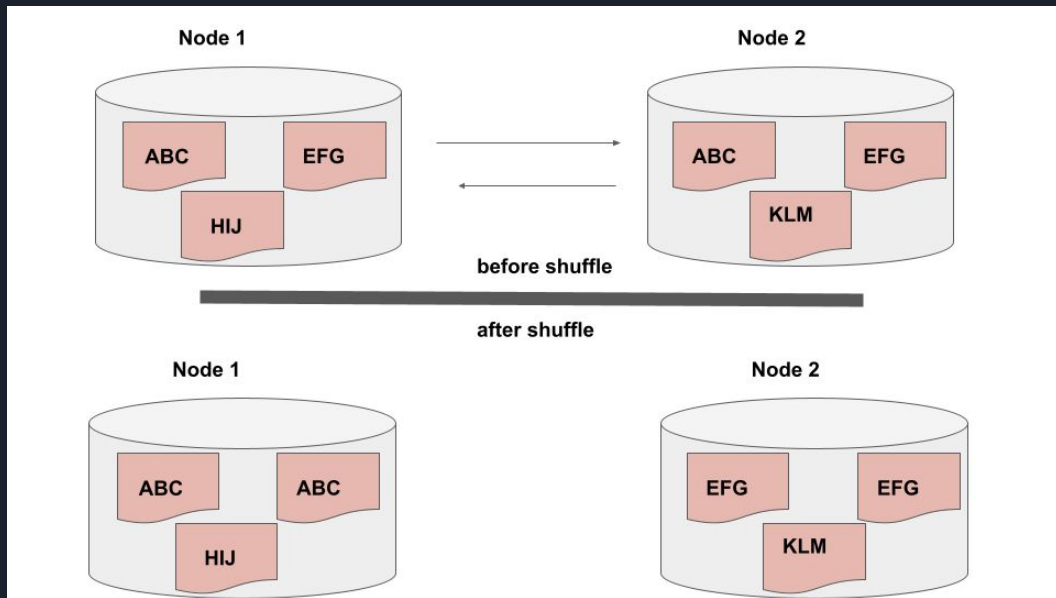
- count, collect, top, take, write/save...



# Data Processing

## Apache Spark

- Data Shuffling



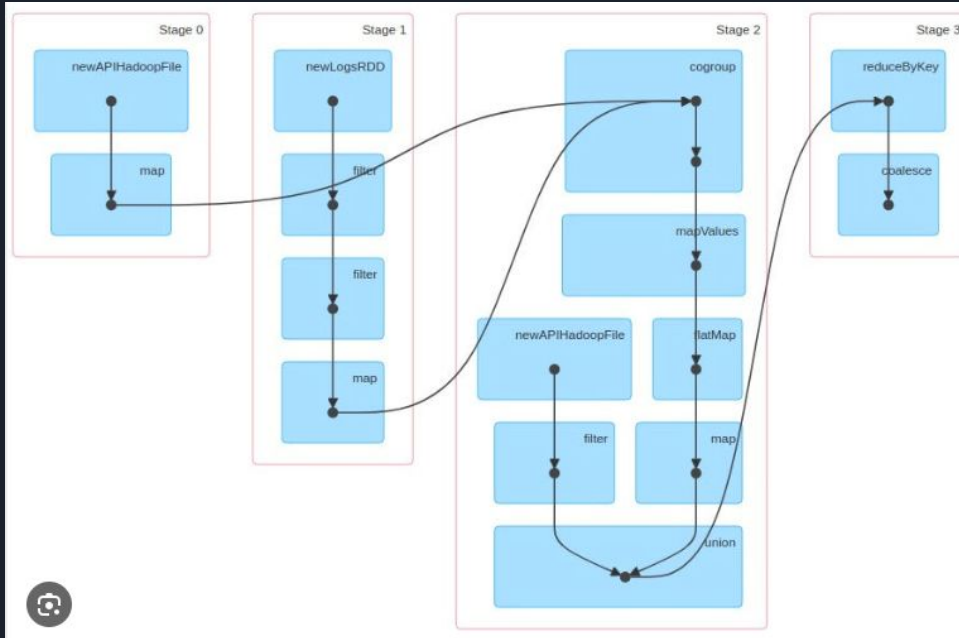
- ABC and EFG in different / same nodes
- Wide transformations
- Partitioning
- Cluster configuration
- Data sizing
- `spark.sql.shuffle.partitions`
  - 200 (default)



# Data Processing

Apache Spark

- Spark DAG Example



- Stages
- Tasks



# Data Processing

Apache Spark

Spark has many properties

- `spark.executor.memory`
- `spark.executor.cores`
- `spark.executor.memoryOverhead`
- `spark.driver.memory`
- `spark.driver.cores`
- `spark.sql.shuffle.partitions`
- ...

To get/set properties: `spark.conf.get("property")` & `spark.conf.set("property", value)`

Spark Docs: <https://spark.apache.org/docs/3.5.1/configuration.html>



**break.**

# Apache Spark Common Issues





# Data Processing

Apache Spark

- Common issues
  - **Having big files / few partitions** - Less executors will be used, rest of executors will be idle (less parallelism = less performance)
  - **Having too many small files** - Requires more network communication for small files and increase data shuffling across the workers
  - **Wrong partition logic in tables** - Avoid columns with high cardinality, choose columns that can be used in filters and aggregations
  - **Skewed data** - Data distribution is not correct on partitions (adjust partition logic / salting technique)
  - **Data Shuffle errors** - Joins/group by can cause wrong data distribution across partition / adjust shuffle partitions size
  - **OutOfMemory** - JVM error, driver or executor run out of memory
  - Performing big transformations that requires data shuffling, using not optimal configuration or processing big amount of data with small resources
  - UDFs performance



# Data Processing

Apache Spark

## HANDS-ON

<https://github.com/lucprosa/dataeng-basic-course/tree/main>

spark/misc/read_from_api.ipynb	Reading from API
spark/misc/etl_program.ipynb	ETL program template
spark/misc/word_count.ipynb	Example of using RDDs



# Data Processing

Lakehouse, Medallion Architecture

- **Lakehouse** = Data Warehouse + Data Lake
  - *Query Engines, data governance, ACID transactions, all type of data, etc*
- **Query Engines** - Query data from databases and data lakes, provides many features like ACID transactions, time-travel, better read/write performance, etc (delta.io, Iceberg, Hudi)
- **Medallion Architecture** - Design pattern to organize data in the data lake
  - How to organize data in a data lake?
  - Transient, Staging, Bronze, Silver (enriched), Gold (curated)

Reference:

Lakehouse - <https://www.databricks.com/glossary/data-lakehouse>

Delta Lake - <https://delta.io/>

Medallion - <https://www.databricks.com/glossary/medallion-architecture>

# Data Processing

ETL flow



- ETL / Spark jobs
  - Ingestion -> Ingest from external data sources and write into bronze/raw layer
  - Cleansing -> Read from bronze, apply transformations and write into silver layer
  - Enrich -> Read from silver, apply business logic/aggregations and write into gold layer



# Data Processing

ETL flow

## 1. import libraries

```
from pyspark.sql import DataFrame
```

## 2. Read the data sources

```
# reading from a table
df = spark.table("sales_db.sales")

# reading from a parquet file
df_1 = spark.read.parquet("/mnt/sales_db/sales")

# reading from a csv file
df_2 = spark.read.format("csv").load("/mnt/sales_db/sales")
```

## 3. Apply transformations

```
# rdd 1
df = df.filter("col1 == 'a'")

# rdd 2
df = df.limit(100)

# rdd3
df = df.join(df_1, "col1").select("col1", "col2", "col3")

# rdd4
output = df.union(df_2)
```

## 4. Write into target

```
output.write.format("delta").saveAsTable("sales_db.new_table")
```



# Data Processing

Data Orchestrator

How to run / orchestrate Spark jobs?

- Apache Airflow - <https://airflow.apache.org/>
- Prefect - <https://www.prefect.io/>
- Azkaban - <https://azkaban.github.io/>
- Azure Data Factory
- Databricks Workflow
- Cron Linux




# Data Processing

Apache Spark

## SPARK CHALLENGES

<https://github.com/lucprosa/dataeng-basic-course/tree/main/spark/challenges>



# Data Engineering

- Data Processing
- Batching
- Spark

Data Engineering Course  
Lucas Rosa  
2024